Multi-Seed Validation — Is the GRU Result Real?

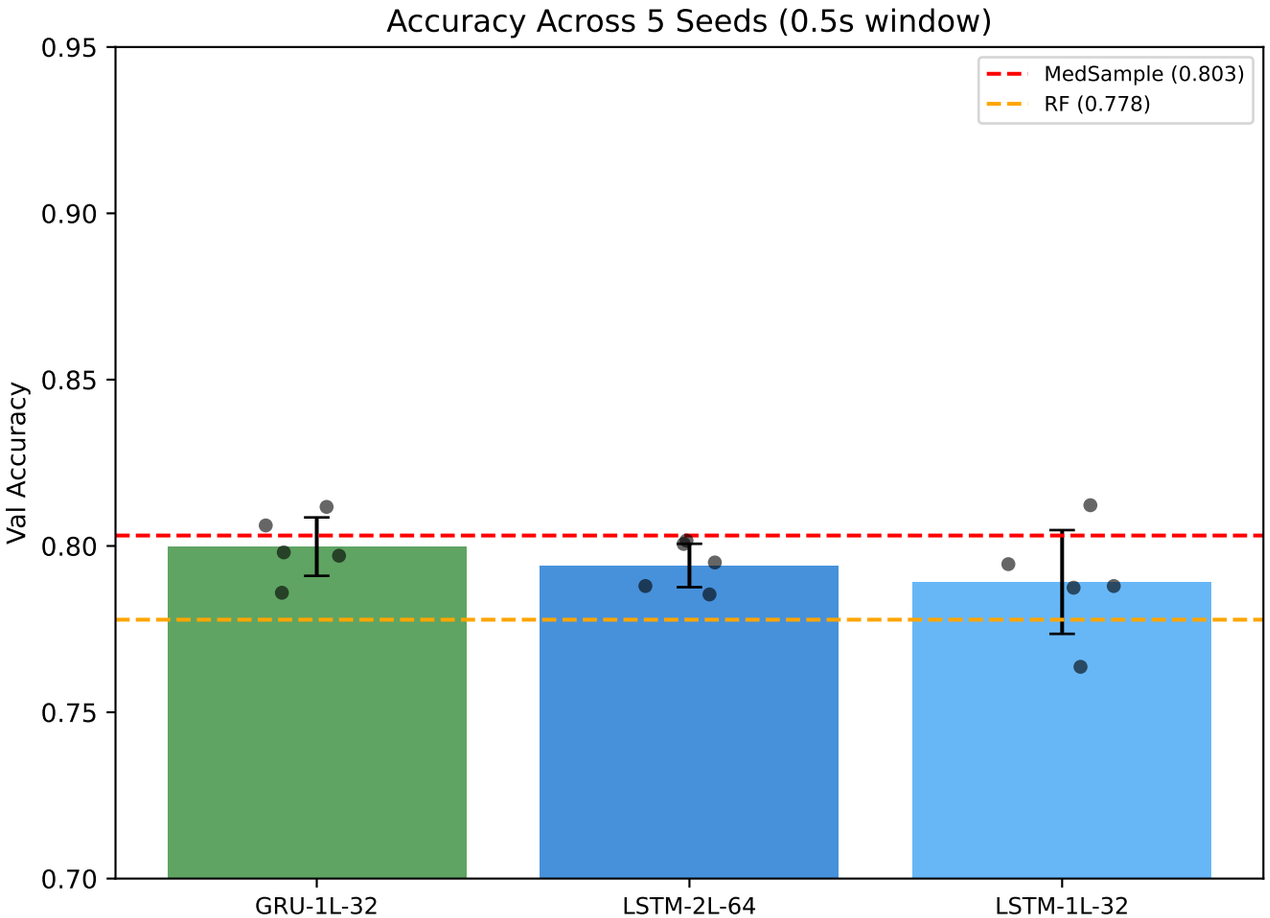

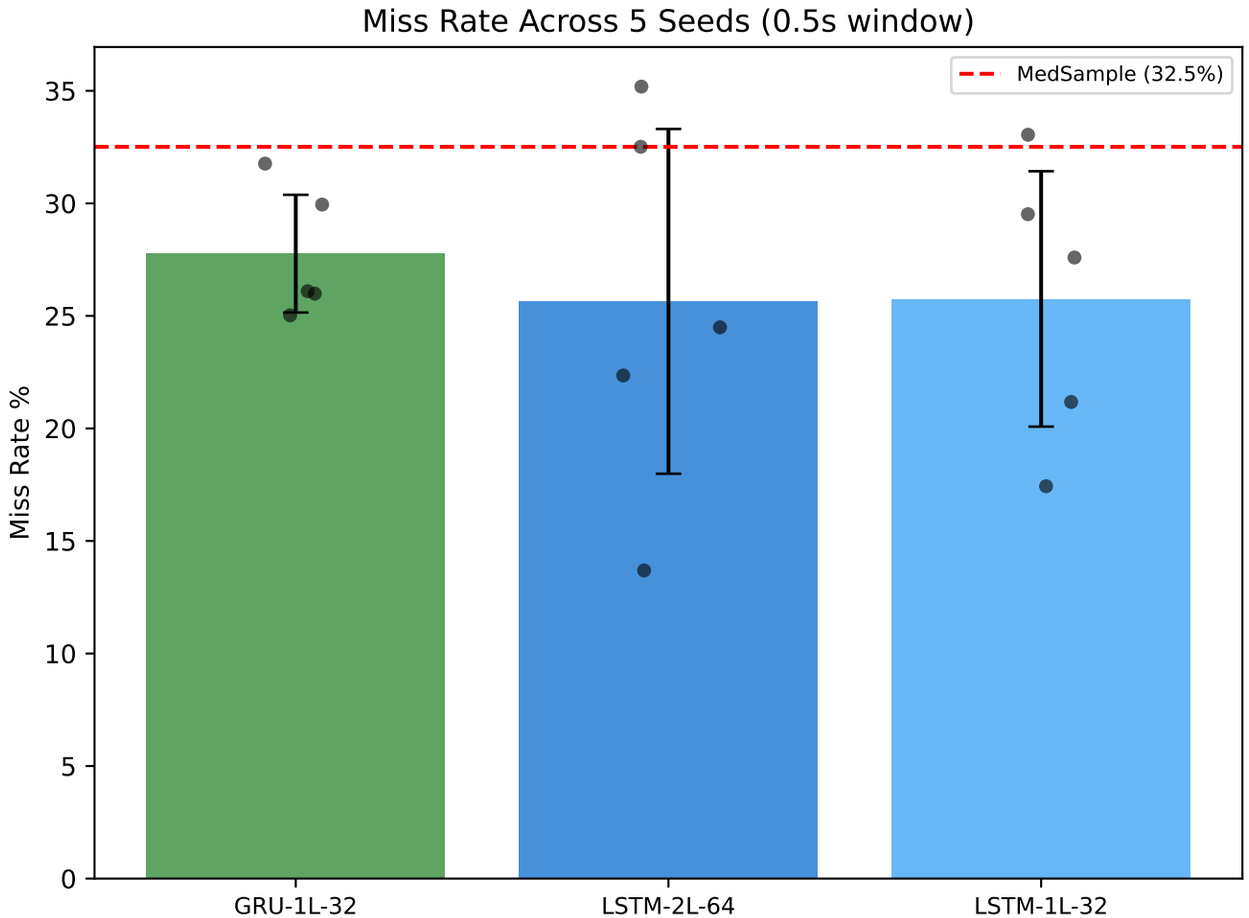

```
======================================================================
MULTI-SEED VALIDATION SUMMARY (0.5s window, DETRAC)
======================================================================

  GRU-1L-32:
    Acc:  0.800 ± 0.009 [0.786 – 0.812]
    F1:   0.773 ± 0.013
    AUC:  0.862 ± 0.015
    Miss: 0.278 ± 0.026

  LSTM-2L-64:
    Acc:  0.794 ± 0.007 [0.785 – 0.802]
    F1:   0.772 ± 0.020
    AUC:  0.854 ± 0.019
    Miss: 0.256 ± 0.077

  LSTM-1L-32:
    Acc:  0.789 ± 0.016 [0.764 – 0.812]
    F1:   0.769 ± 0.020
    AUC:  0.856 ± 0.020
    Miss: 0.258 ± 0.057

  Baselines (deterministic):
    MedSample: Acc=0.803 Miss=32.5%
    RF:        Acc=0.778 Miss=24.2%

======================================================================
VERDICT: GRU does NOT reliably beat MedSample (0.800±0.009 vs 0.803)
```